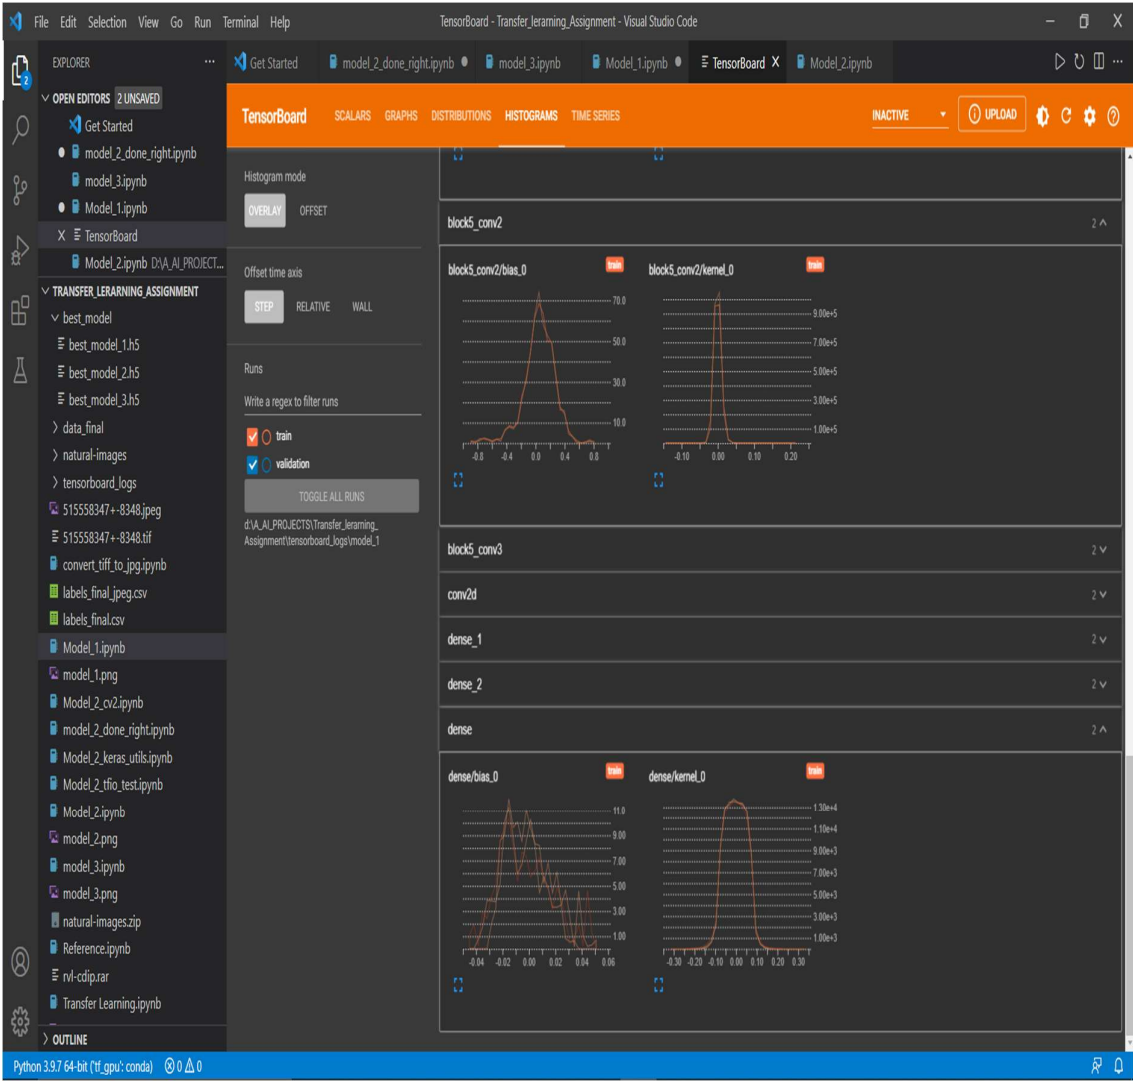
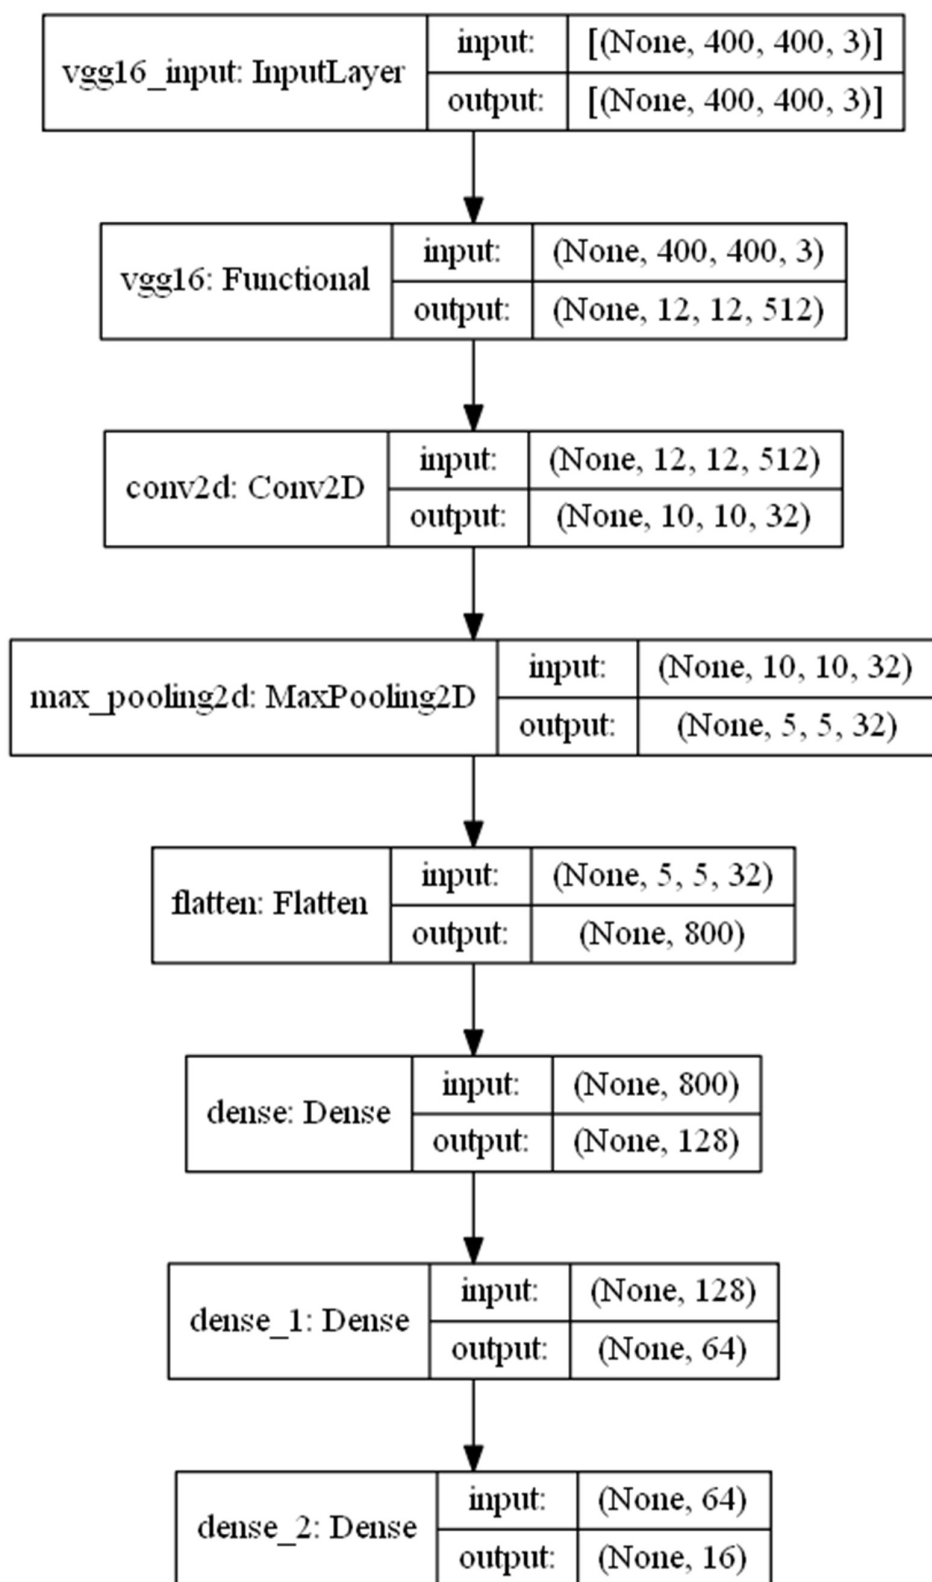
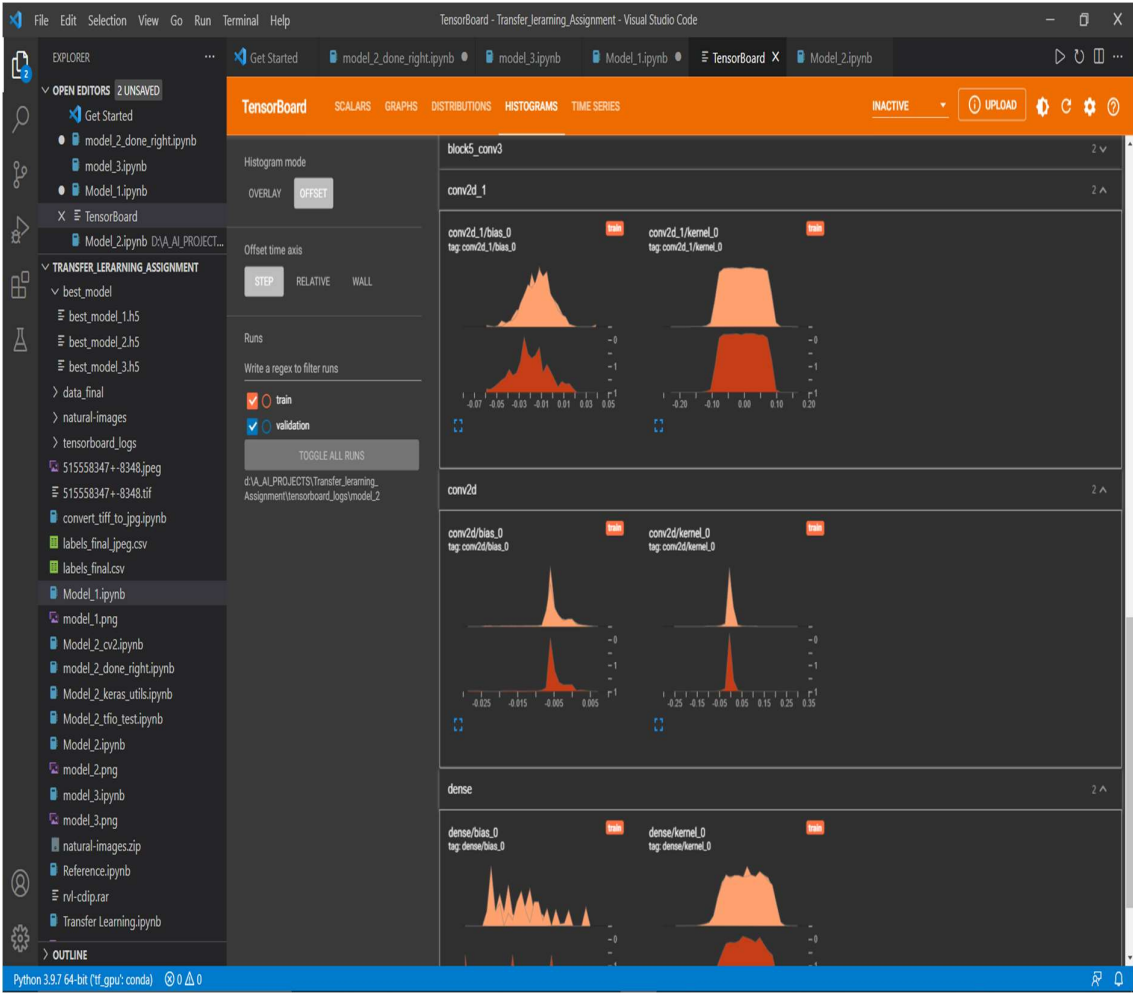


Model 1.





Model 2.



|                         |         |                       |
|-------------------------|---------|-----------------------|
| vgg16_input: InputLayer | input:  | [(None, 400, 400, 3)] |
|                         | output: | [(None, 400, 400, 3)] |



|                   |         |                     |
|-------------------|---------|---------------------|
| vgg16: Functional | input:  | (None, 400, 400, 3) |
|                   | output: | (None, 12, 12, 512) |



|                |         |                     |
|----------------|---------|---------------------|
| conv2d: Conv2D | input:  | (None, 12, 12, 512) |
|                | output: | (None, 1, 1, 512)   |



|                  |         |                   |
|------------------|---------|-------------------|
| conv2d_1: Conv2D | input:  | (None, 1, 1, 512) |
|                  | output: | (None, 1, 1, 256) |

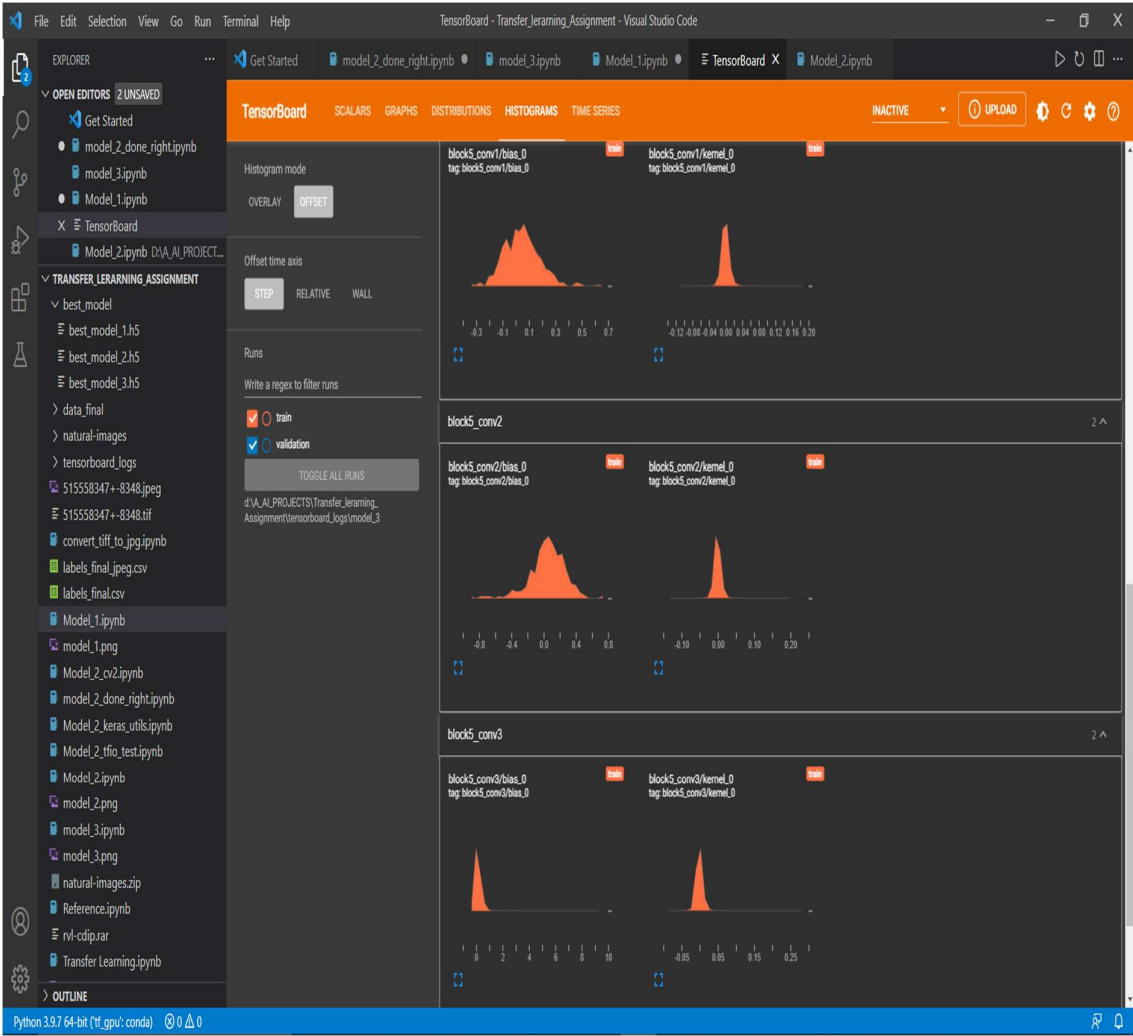


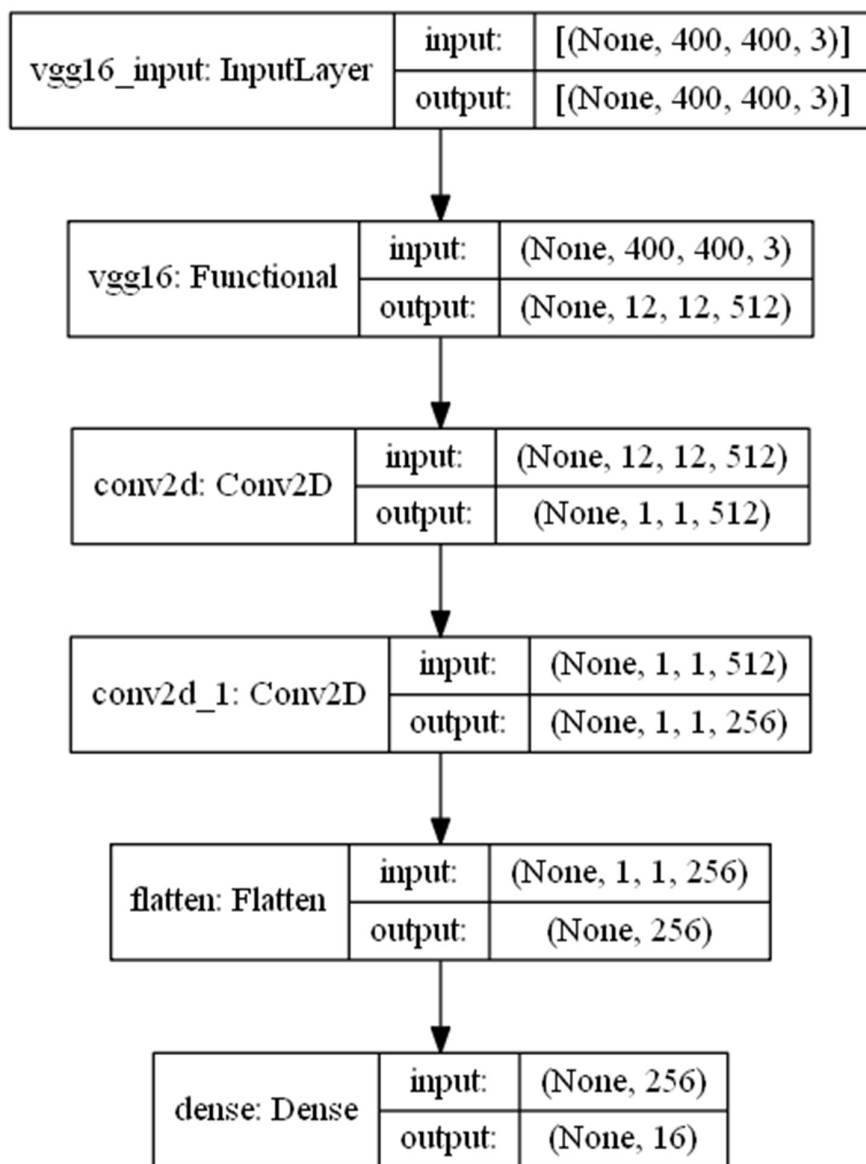
|                  |         |                   |
|------------------|---------|-------------------|
| flatten: Flatten | input:  | (None, 1, 1, 256) |
|                  | output: | (None, 256)       |



|              |         |             |
|--------------|---------|-------------|
| dense: Dense | input:  | (None, 256) |
|              | output: | (None, 16)  |

Model 3.





#### Observations.

- Model 1 and 2 performed similarly, both gained an accuracy of around 55% in the first epoch itself.
- Model 2 was a bit faster to train than Model 1.
- Model 3 took more time to train and was less accurate.
- Using ImageDataGenerator from Keras results in data-starvation to the GPU.  
Using tf.data module cuts time in half.